

CO 327 – Operating Systems

Multiprocessing

DEPARTMENT OF COMPUTER ENGINEERING

The simple servers you write in introductory *Socket Programming* exercises only read or write to a single socket at a time. So, how is a web server like *Apache* able to serve thousands of clients simultaneously? To do that we need to be able to perform reads and writes to sockets *concurrently*. There are three concurrency techniques provided by UNIX:

1. Multiprocessing
2. Multithreading
3. Non blocking I/O

This lab introduces the first technique.

1. Processes

Running a program on UNIX creates a *process*. The kernel creates a *process table* entry to record the resources used by the process, such as memory and open file descriptors. Each process is identified by a unique *process identifier* (PID).

Exercise 1: Use the following commands to view details of the processing running on your system. Note the PIDs.

- i. *top* shows you details of active processes. The processes are sorted by CPU usage by default. Sort them by memory usage.
- ii. Run *ps* with the following options: *-a*, *-x*, *-u*, *-w*. What is the name of the process with PID 1?

1.1. Creating a new process

To programmatically create a process UNIX provides the *fork()* system call. When *fork()* is called it creates a new process and returns twice, once in the *parent process* and once in the new *child process*. To check whether you are in the parent or child, you check the return value of *fork()*.

```
int main(void)
{
    int pid;
    pid = fork();
    if (pid < 0)
    {
        perror("fork");
        exit(1);
    }
    if (pid == 0)
        puts("This is the child process");
    else
        puts("This is the parent process");

    return 0;
}
```

This is the standard pattern that is used when calling *fork*.

`fork()` is easy, since it can only return three things:

- 0 If it returns 0, you are the child process. You can get the parent's PID by calling `getppid()`. Of course, you can get your own PID by calling `getpid()`.
- 1: If it returns -1, something went wrong, and no child was created. Use `perror()` to see what happened.
- else: Any other value returned by `fork()` means that you're the parent and the value returned is the PID of your child. This is the only way to get the PID of your child, since there is no `getcpid()` call. Why?

Exercise 2:

1. In what order are the messages from parent and child printed? Is the order *always* the same?
2. How many children will the following program spawn? Draw a diagram illustrating the parent-child relationships between processes.

```
int main(void)
{
    for (int i=0; i<3; i++)
        fork();
}
```

1.2. Waiting for children

The system call `wait()` lets the parent process wait until the child process has exited. For example, a shell must wait until a command a user has run completes before prompting the user for the next command.

Exercise 3: Modify the program in section 1.1 so that the parent always prints its message *after the child*. Refer to `man 2 wait` for details.

1.3. Replacing the process image

In certain cases we would like to execute another program within a process. For example, a shell must create a new process and then run an external program within that process. This is made possible by the `exec()` system call. Doing an `exec` replaces the current process image in memory with a new program. Therefore a call to `exec` does not return.

```
int main(char argc, char **argv)
{
    execl("/bin/ls", "-l", argv[1], NULL);
    puts("Program ls has terminated");
}
```

This example is using the `execl()` variation provided by the standard library. See `man 3 exec` for details.

Exercise 4:

1. Compile and run the above code giving it a path as an argument. How many times is the message “*Program ls has terminated*” printed?
2. Write a very simple shell that repeatedly prompts the user for a command and runs it with any arguments given. Make sure your shell waits until the command has completed before prompting the user for the next command.

2. Multiprocess servers

We can now apply these techniques to build servers that concurrently handle multiple client requests using multiple server processes. A socket is set up to `listen()` for client connection requests in the same way as iterative servers. When a new client request arrives `accept()` returns a new socket connected to the client.

The main loop of a multiprocess server is where the difference lies. Instead of handling the request itself, the server spawns a child process to handle the client while the parent process continues to `listen()` for new connections. In this way the server is able to handle multiple clients concurrently.

```
listen(sockfd,5);
clilen = sizeof(cli_addr);
while (1)
{
    /* New socket descriptor is returned each time a client connects*/
    newsockfd = accept(sockfd, (struct sockaddr *) &cli_addr, &clilen);
    if (newsockfd < 0)
    {
        perror("ERROR on accept");
        exit(1);
    }
    pid = fork();
    if (pid < 0)
    {
        perror("ERROR on fork");
        exit(1);
    }
    if (pid == 0)
    {
        /* In child process which the handles client connection */
        close(sockfd);
        handle_client(newsockfd);
        exit(0);
    }
    else
        /* In parent process which continues to listen for new clients */
        close(newsockfd);
}
```

Exercise 5:

1. Open three terminals and run the server in one. Use `nc()` to connect as two clients concurrently on port 12345. Type some text in both clients and examine the client and server outputs.
2. Suppose we modify the server parent process to call `wait()` on the last line above (**highlighted**) to wait until the child serving a client terminates. What would happen?
3. What happens if you terminate the the server while a client is connected, and then try to restart it? (Resolving this issue requires a *signal handler*.)
4. Modify this server to do the following: The client sends the path to a file whose contents the server will send back to the client (if the file exists.) Verify that your new server can handle multiple concurrent connections by using `nc()`. Can two concurrent clients request the same file?